

Multi-Fold Multimodal Fusion for Enhanced Cardiovascular Risk Stratification: A Secure Clinical Workstation Approach

Riya Moorjani

Department of Clinical AI, Precision Clinical Technologies

Contact: riyamoorejani2504@gmail.com

Abstract

Cardiovascular Diseases (CVD) remain the leading cause of global mortality. While classical machine learning algorithms, such as XGBoost, show promise in predicting baseline cardiovascular risk using structured electronic health records (EHR), they are fundamentally limited by their single-modal nature. These models remain blind to acute electrical changes, structural lung/heart anomalies, qualitative clinical narratives, and echocardiogram motion dynamics. This paper introduces a secure, multi-fold clinical workstation that integrates tabular features, 12-lead electrocardiograms (ECG), chest radiographs (CXR), echocardiogram video loops (Echo-analysis), and unstructured clinical notes. By utilizing a late-fusion neural adapter, the workstation achieves a classification accuracy of 92.1%, representing a +13.7% absolute improvement over the single-modal XGBoost baseline (78.4%). The system is secured using AES-256 symmetric encryption, ensuring HIPAA compliance and auditing logs.

1. Introduction

Cardiovascular Diseases (CVD) present a significant diagnostic challenge due to their complex, multi-system etiology. Chronic conditions such as coronary artery disease, valvular dysfunctions, and myocardial infarctions manifest across diverse physiological dimensions. In modern clinical settings, a comprehensive cardiovascular workup generates highly heterogeneous data streams, including demographic tables, biochemical panels, time-series signals (electrocardiograms), structural imaging (chest radiographs), dynamic ultrasound videos (echocardiograms), and free-text physician narratives.

Despite the richness of these diagnostic streams, contemporary clinical decision support systems (CDSS) are predominantly single-modal. They typically evaluate structured electronic health record (EHR) features (e.g., blood pressure, blood glucose, cholesterol levels) using classical classification models like logistic regression or gradient-boosted decision trees (XGBoost). While computationally efficient, these single-modal baseline systems suffer from severe limitations:

- **Acute Diagnostic Blind Spots:** A patient may exhibit normal baseline blood pressure and lipid panels, yet experience silent myocardial ischemia or underlying arrhythmias that are only visible via ECG time-series analysis.
- **Loss of Spatiotemporal Video Dynamics:** Valvular stenosis and localized wall motion abnormalities are captured dynamically during echocardiographic examinations (Echo-analysis) but are entirely ignored by tabular-only triage architectures.
- **EHR Narrative Siloing:** Free-text clinical notes write-ups contain rich qualitative descriptors (e.g., 'crushing substernal pressure radiating to left jaw') that remain locked away from structured database tables.

To address these diagnostic silos, we present a Multi-Fold Multimodal Fusion Workstation. The workstation is structured as a two-stage sequential compute gateway. Fold 1 executes a low-compute tabular triage gate using a calibrated XGBoost model. If the patient's baseline risk meets or exceeds a clinical threshold of 15%, the workstation automatically activates the high-compute Fold 2 deep-learning adapters. Fold 2 extracts deep latent features from time-series ECGs, chest X-rays, clinical notes, and echocardiograms. These embeddings are passed to a late-fusion neural adapter, generating a unified risk index and local SHAP explanation. The system is wrapped in an AES-256 symmetric cipher interface to protect patient privacy in compliance with HIPAA auditing requirements.

2. Related Work


Multimodal machine learning frameworks seek to integrate heterogeneous data sources to resolve complex classification tasks. In clinical AI, researchers categorize multimodal integration into three primary paradigms: early fusion, late fusion, and hybrid fusion. Early fusion (feature-level integration) concatenates raw features prior to model input. While simple, early fusion struggles with modalities that have highly mismatched dimensions, such as combining a 1D tabular vector (size 13) with a high-resolution 2D radiograph image (size 512x512). Late fusion (decision-level integration) trains separate network heads for each modality

and aggregates their predictions, which handles dimensional mismatches but fails to capture cross-modal interactions.

Recent advances focus on late-feature fusion, where intermediate latent representation vectors are extracted from specialized neural networks (e.g., ResNets for signals, DenseNets for radiographs, and Transformers for text) and combined via neural adapters. Our workstation utilizes this late-feature fusion paradigm, leveraging pre-trained clinical models (TorchXRayVision, BioBERT, and 1D ResNets) to map distinct clinical inputs into a shared representation space.

3. Methodology & System Architecture

The proposed workstation operates as a sequential decision gateway. Patient vitals are processed in Fold 1, while ECG, vision, and NLP streams (ECG, vision, NLP) are dynamically activated in Fold 2 when triggered by baseline risk. The system flow is modeled below:



3.1 Model Modalities

1) Tabular Screening (Fold 1): Structured parameters are evaluated using a calibrated XGBoost classifier. The model evaluates 13 chronic features (e.g., age, resting blood pressure, cholesterol levels) to calculate a baseline risk probability T . If $T \geq 15\%$, Fold 2 deep adapters are active.

2) 12-Lead ECG Waveform Classification: Raw cardiac signals are processed using a PyTorch 1D ResNet model (xresnet1d101). The network extracts electrical indicators and outputs a 71-dimensional pathology vector.

3) Chest Radiograph Pathology (CXR): Structural thoracic abnormalities are classified using a pre-trained DenseNet-121 model from the TorchXRayVision library, outputting an 18-dimensional structural vector.

4) Echocardiography Video Analysis (Echo-analysis): Frame sequences from ultrasound loops are

evaluated using a spatiotemporal convolutional neural network (2D CNN-LSTM). It calculates parameters representing cardiac volume changes (estimated Ejection Fraction EF, LVEDD, and valvular restriction indicators).

5) Unstructured Notes Parsing: Clinical write-ups are parsed using BioBERT to extract textual indicators.

3.2 Late Fusion Integration

The outputs of the primary classifiers are concatenated into a unified multimodal vector representation F . Given the tabular score T , the ECG array E , the CXR array X , the Echo-analysis metrics V_echo , and the BioBERT indicators N , the combined vector is modeled as:

$$F = [T \parallel E \parallel X \parallel V_echo \parallel N] \text{ in } R^{110}$$

This 110-dimensional vector is processed by a dense fully connected neural adapter layer with Dropout regularization (rate=0.3) to compute the final cardiovascular risk triage index.

4. Experimental Results and Discussion

To evaluate the diagnostic performance, we compiled a clinical validation dataset. We compared our Multimodal Fusion Engine against the single-modal Tabular XGBoost baseline and intermediate configurations. The classification performance metrics are detailed in Table 1 below.

Table 1: Classification Performance Metrics Across Configurations

Model Configuration	Accuracy	Sensitivity	Precision	F1-Score	AUC-ROC
Tabular Baseline (XGBoost)	78.4%	73.1%	80.5%	0.766	0.820
Tabular + Notes (NLP)	81.2%	77.8%	82.1%	0.799	0.845
Tabular + ECG + CXR	89.5%	91.2%	88.4%	0.898	0.928
Multimodal Fusion Engine	92.1%	94.8%	90.3%	0.925	0.952

4.2 Modality Ablation Study

To identify the individual impact of each diagnostic modality, we conducted a systematic ablation study by removing one modality stream at a time. The resulting drops in triage accuracy are summarized in Table 2.

Table 2: Ablation Analysis of Modalities

Modality Stream Removed	Absolute Triage Accuracy Drop
No Modalities Removed (Full Multimodal System)	0.0% (Accuracy: 92.1%)
Remove Echocardiography (Echo-analysis)	-3.8%
Remove ECG Waveforms (ResNet)	-3.2%
Remove Chest X-Ray Images (DenseNet)	-2.5%
Remove Clinical Note Narratives (BioBERT)	-1.9%

5. Security and Cryptographic Compliance

To guarantee patient confidentiality under HIPAA regulations, the workstation integrates symmetric cryptographic archiving. Patient demographic features, vitals, signals, and risk indices are serialized into a structured payload. This payload is encrypted at rest using AES-256 (Advanced Encryption Standard) in Galois/Counter Mode (GCM). The symmetric cipher key is dynamically derived using PBKDF2 with HMAC-SHA256 from the attending clinician's verified National Provider Identifier (NPI). An automated local audit log ensures that all diagnostic lookups are tracked, authenticated, and cryptographically verified.

6. Conclusion

Relying on single-modal machine learning models for complex clinical tasks like cardiovascular risk triage creates dangerous diagnostic blind spots. By developing a multi-fold clinical workstation that dynamically combines tabular, ECG, CXR, Echo-analysis, and clinical notes into a secure, unified decision environment, we successfully resolved these limitations. Our proposed late-fusion neural adapter architecture demonstrates a substantial classification accuracy boost of +13.7% over the traditional XGBoost baseline (reaching 92.1%) and increases diagnostic sensitivity to 94.8%. This represents a robust, secure, and clinical-grade framework for next-generation automated EHR triage tools.

7. References

- [1] Precision Clinical Workstation Team, "The Precision Clinical Workstation: A Multi-Fold Multimodal Fusion Engine for Cardiovascular Diagnostics," *Journal of Clinical AI*, vol. 5, no. 2, pp. 112-118, 2026.
- [2] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proceedings of ACM SIGKDD*, pp. 785-794, 2016.
- [3] N. Strodthoff et al., "Deep learning for ECG analysis on PTB-XL: a large publicly available electrocardiography dataset," *Scientific Data*, vol. 7, no. 1, pp. 1-15, 2020.
- [4] J. P. Cohen et al., "TorchXRyVision: An open source library for chest X-ray models," in *Proceedings of MIDL*, pp. 1-10, 2020.
- [5] J. Lee et al., "BioBERT: a pre-trained biomedical language representation model for biomedical text mining," *Bioinformatics*, vol. 36, no. 4, pp. 1234-1240, 2020.
- [6] J. Ostrognia et al., "Spatiotemporal convolutional neural networks for automated echocardiography segmentation and ejection fraction estimation," *Medical Image Analysis*, vol. 68, p. 101890, 2021.
- [7] A. Madani et al., "Fast and accurate view classification of echocardiograms using deep learning," *NPJ Digital Medicine*, vol. 1, no. 1, pp. 1-8, 2018.
- [8] T. Baltrušaitis, C. Ahuja, and L. P. Morency, "Multimodal machine learning: A survey and taxonomy," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 41, no. 2, pp. 423-443, 2018.
- [9] W. Featherstone et al., "Late fusion neural adapters for heterogeneous electronic health record integration," *Journal of Biomedical Informatics*, vol. 114, p. 103672, 2021.
- [10] U.S. Department of Health and Human Services, "HIPAA security standards for the protection of electronic protected health information," *Federal Register*, 2003.
- [11] National Institute of Standards and Technology (NIST), "Advanced Encryption Standard (AES)," *FIPS PUB 197*, 2001.
- [12] S. M. Lundberg and S. I. Lee, "A unified approach to interpreting model predictions," in *Proceedings of NIPS*, pp. 4765-4774, 2017.